

Flickr8K dataset:

Flickr8K_Dataset: https://drive.google.com/file/d/1u3oqx36XApnAykFDB6EEWUIfd_CxRQQ9/view

(Contains around 8K images (8091))

Flickr8K_text: <https://drive.google.com/file/d/1qcRy3WpQv4dGtu65gETtYLWxDpBrRtx1/view>

(Contains captions

- 1) Flickr8k.lemma.token.txt : 5 captions per image for all 8K images.
- 2) Flickr_8k.trainImages.txt : List of 6K training images set.
- 3) Flickr_8k.testImages.txt : List of 1K test images set.
- 4) Flickr_8k.devImages.txt : List of 1K dev/validation images set.

)

In `main.py`:

Data Preprocessing and Cleaning

- 1) **Load caption data** using `load_doc()` by reading the text file as a single string.
- 2) **Parse raw captions** using `all_img_captions()`, grouping multiple captions per image into a dictionary.
- 3) **Clean caption** with `cleaning_text()` by removing punctuation, converting to lowercase, filtering out short and non-alphabetic words.
- 4) **Build a vocabulary** of unique words from the cleaned captions using `text_vocabulary()`.
- 5) **Save processed captions** to a new file (`descriptions.txt`) in a structured format.

We have cleaned 5 captions per image saved in “**descriptions.txt**” for all 8K images.

Image Feature Extraction Using Pretrained CNN (Xception)

- 6) **Download pretrained Xception weights** from TensorFlow’s model zoo. **Initializes the Xception model** (excluding top classification layer) for feature extraction using `include_top=False`.
- 7) **Initialize the Xception model** (excluding top classification layer) for feature extraction using `include_top=False`.
- 8) **Define `extract_features()`** to process each image:
 - Loads and resizes the image to 299×299 (Xception input size),
 - Normalizes pixel values,
 - Passes it through the Xception model to extract a **2048-dimensional feature vector**.
- 9) **Save extracted image features** as a dictionary using `pickle` (`features.p`) for future use.

We have extracted features from all 8K images and saved it in “**features.p**”

Preparing Training Data

Let’s start working with training dataset:

- 10) **`load_photos()`**: Load a list of training image filenames from a text file and filter only those images that exist in the training dataset directory.
- 11) **`load_clean_descriptions()`**: Load preprocessed (cleaned) image captions from a file, keeping only those that correspond to the training images, and wraps each caption with `<start>` and `<end>` tokens.
- 12) **`load_features()`**: Load the saved image features (`features.p`) and filter them to include only features corresponding to the training images.

Tokenizer

Let's make tokenizer:

- 13) `dict_to_list()`: Convert the dictionary of image-to-captions into a flat list of all captions.
`create_tokenizer()`: Use Keras's `Tokenizer` to fit on the list of captions and map words to unique integer indices.
- 14) `create_tokenizer()`: Use Keras's `Tokenizer` to fit on the list of captions and map words to unique integer indices.
- 15) **Save the tokenizer** using `pickle` for future use during training and inference.

Conversion of captions into training sequences

- 16) Calculate `vocab_size`: the total number of unique words in the dataset and compute the maximum length of all training captions.
- 17) `create_sequences()`
For each caption:
 - Converts it to a sequence of integers using the tokenizer.
 - Generates multiple input-output pairs where:
 - Input = image feature + partial caption (padded),
 - Output = next word (one-hot encoded).
 - Example: `'a dog runs'` →
 - Input: `['a']` → Output: `'dog'`
 - Input: `['a', 'dog']` → Output: `'runs'`
- 18) `data_generator()`
Creates a `tf.data.Dataset` generator that:
 - Yields batches of `{image feature, partial caption} → next word` pairs,
 - Matches the model's expected input-output shape using `tf.TensorSpec`.

Model Architecture and Training

- 19) Model Definition:
 - Image input branch:**
 - Accepts 2048-dimensional image feature vectors (from CNN),
 - Applies dropout and reduces to 256 features using a dense layer.
 - Caption input branch:**
 - Accepts tokenized caption sequences (max length),
 - Embeds words into 256-dimensional vectors,
 - Processes the sequence using an LSTM to capture context.
 - Merging branches:**
 - Combines image features and LSTM output using `add()`,
 - Passes through dense layers to predict the next word using a `softmax` layer.
 - Compiles the model** with `categorical_crossentropy` loss and `adam` optimizer.
- 20) Model training:
 - Runs training for 10 epochs, After each epoch, saves the model to a new `.h5` file inside the `models/` directory.

In `evaluate.py`:

- 1) **Defines helper functions** to:
 - Load image filenames, captions, and pre-extracted image features,
 - Convert predicted word indices back to words,
 - Generate captions for images using the trained model.
- 2) **Defines a multimodal model** architecture that combines CNN image features and LSTM-generated text context, to predict the next word in the caption sequence.
- 3) **Loads pre-trained tokenizer and sets vocab size and max caption length.**
- 4) **Loads test image features and cleaned test captions.**
- 5) **Loads model weights from a `.h5` file**, compiles the model, and prepares it for evaluation.
- 6) **Generates captions for test images** and compares them to ground-truth captions using **BLEU scores (1–4)**.
- 7) **Prints BLEU scores** as a quantitative evaluation of the model's captioning accuracy.

Results:

The model was trained with a vocabulary size of **7,577** and a maximum caption length of **32**. It contained 5,002,649 trainable parameters. Training was conducted for 10 epochs, with the loss steadily decreasing from **5.03 to 2.92**, indicating good convergence.

7700/7700	282s 37ms/step - loss: 5.0365
7700/7700	266s 34ms/step - loss: 3.8218
7700/7700	278s 36ms/step - loss: 3.4940
7700/7700	276s 36ms/step - loss: 3.3104
7700/7700	279s 36ms/step - loss: 3.1918
7700/7700	280s 36ms/step - loss: 3.1070
7700/7700	280s 36ms/step - loss: 3.0502
7700/7700	281s 36ms/step - loss: 2.9944
7700/7700	283s 37ms/step - loss: 2.9602
7700/7700	280s 36ms/step - loss: 2.9220

BLEU Score of validation set for 10 epochs, demonstrate general improvement and then stabilization.

Epoch 0:	BLEU-1=0.3624, BLEU-2=0.1912, BLEU-3=0.1199, BLEU-4=0.0487
Epoch 1:	BLEU-1=0.3458, BLEU-2=0.1856, BLEU-3=0.1206, BLEU-4=0.0489
Epoch 2:	BLEU-1=0.2871, BLEU-2=0.1571, BLEU-3=0.1060, BLEU-4=0.0415
Epoch 3:	BLEU-1=0.3150, BLEU-2=0.1680, BLEU-3=0.1095, BLEU-4=0.0422
Epoch 4:	BLEU-1=0.3108, BLEU-2=0.1649, BLEU-3=0.1067, BLEU-4=0.0408
Epoch 5:	BLEU-1=0.3431, BLEU-2=0.1860, BLEU-3=0.1265, BLEU-4=0.0539
Epoch 6:	BLEU-1=0.3546, BLEU-2=0.1904, BLEU-3=0.1281, BLEU-4=0.0549
Epoch 7:	BLEU-1=0.3345, BLEU-2=0.1798, BLEU-3=0.1174, BLEU-4=0.0452
Epoch 8:	BLEU-1=0.3589, BLEU-2=0.1953, BLEU-3=0.1332, BLEU-4=0.0574
Epoch 9:	BLEU-1=0.3451, BLEU-2=0.1871, BLEU-3=0.1269, BLEU-4=0.0535

During validation, the best BLEU-4 score was achieved at epoch 8, with BLEU-4=**0.0574**

On the test set, the model from the last epoch (epoch 9) achieved:

BLEU-1: 0.333130 BLEU-2: 0.181531 BLEU-3: 0.120728 BLEU-4: 0.047832

However, the model from epoch 8 (with best validation BLEU-4) gave improved results on the test set:

BLEU-1: 0.346744 BLEU-2: 0.185889 BLEU-3: 0.122901 BLEU-4: 0.051291

This project successfully demonstrates image captioning by integrating pretrained CNN features with LSTM-based language modeling on the Flickr8k dataset. The model achieved a best BLEU-4 score of **0.0574 on validation** and **0.0513 on test data**